

# Optimistic Concurrency Control for Distributed Transactions

May 31, 2021

## Contents

<b>1</b>	<b>Concepts</b>	<b>2</b>
1.1	<b>TODO</b> Strict serialization . . . . .	2
1.2	<b>TODO</b> Optimistic concurrency control . . . . .	2
1.3	<b>TODO</b> Two-phase commit (2PC) . . . . .	2
<b>2</b>	<b>Protocol</b>	<b>2</b>
2.1	Assumptions . . . . .	2
2.2	Requirements . . . . .	2
2.3	Implementation . . . . .	2
2.3.1	Client . . . . .	3
2.3.2	Coordinator . . . . .	3
2.3.3	Server . . . . .	4

Collection of notes and code for the project of the Distributed Systems 1 course, part of the Master in Computer Science at the Università di Trento (UNITN).

## 1 Concepts

### 1.1 **TODO** Strict serialization

### 1.2 **TODO** Optimistic concurrency control

### 1.3 **TODO** Two-phase commit (2PC)

## 2 Protocol

The goal is to allow **multiple clients** to make **concurrent transactions (TXN)** on a **distributed database**. A TXN is a sequence of R/W operations on one or more items that can be committed or aborted.

### 2.1 Assumptions

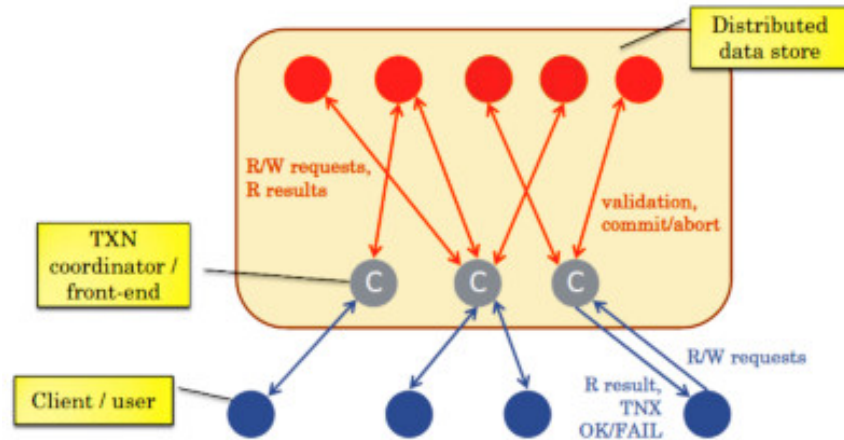
- Links are FIFO and reliable.
- Unicast transmission delays are simulated with small, random intervals.
- Clients don't crash, but coordinators and servers can crash at any time (and recover after a given time).

### 2.2 Requirements

- Crashes should be simulated on critical points of the execution, like validation/commit.
- After a number of transactions takes place, the system allows to check if its state is consistent.

### 2.3 Implementation

The distributed nature requires us to use atomic commitment to guarantee consistency (e.g. 2PC). We assume conflicts are rare, and in doing so we choose to use a private workspace instead of any form of lock (whose adoption in our scenario would be wasteful and non-trivial).



### 2.3.1 Client

Each client tries to start a transaction, one at a time, by sending a `TXN_BEGIN` to a randomly choosed coordinator. Once the coordinator confirms the transaction, the client will send a series of operations (ending with a `TXN_END`) and the TXN will eventually commit or abort.

The client implementation is given, and its operation consist in reading items two-by-two, then for a given value  $x$  subtracting it to the first item and adding it to the second (the sum should stay constant). Clients don't crash.

### 2.3.2 Coordinator

Coordinators are responsible for

- Uniquely identifying the TXNs
- Forwarding and replying to requests from the client
- Changing the values as requested by the client
- Orchestrating concurrent TXNs using distributed commits to ensure strict serializability

TXNs directly run on the private workspaces, and if valid they are later committed by the coordinator. The validation of a workspace is the only phase that would requirer the coordinator to lock (some) items, still for a very short window of time if compared with pessimistic locking.

### 2.3.3 Server

The server stores a part of the distributed data, that consists in a simple collection of key-value pairs (with integer values). A data store server  $S_i$  manages data with keys in the range  $[10i \cdots 10i + 9]$ .